


Iris data

| | sepal length (cm) | sepal width (cm) | petal length (cm) | petal width (cm) | target |
|-------|-------------------|------------------|-------------------|------------------|------------|
| count | 150.000000 | 150.000000 | 150.000000 | 150.000000 | 150.000000 |
| mean | 5.843333 | 3.057333 | 3.758000 | 1.199333 | 1.000000 |
| std | 0.828066 | 0.435866 | 1.765298 | 0.762238 | 0.819232 |
| min | 4.300000 | 2.000000 | 1.000000 | 0.100000 | 0.000000 |
| 25% | 5.100000 | 2.800000 | 1.600000 | 0.300000 | 0.000000 |
| 50% | 5.800000 | 3.000000 | 4.350000 | 1.300000 | 1.000000 |
| 75% | 6.400000 | 3.300000 | 5.100000 | 1.800000 | 2.000000 |
| max | 7.900000 | 4.400000 | 6.900000 | 2.500000 | 2.000000 |



removing unuseful columns

splitting to train data 80% and test data 20% with random selection

scaling data with 0 mean removing the standard deviation


ΣΤΑΤΙΣΤΙΚΑ

NN 4:30:3 with Sigmoid activation

```
NN_Sigmoid(  
  (linear_stack): Sequential(  
    (0): Linear(in_features=4, out_features=30, bias=True)  
    (1): Sigmoid()  
    (2): Linear(in_features=30, out_features=3, bias=True)  
  )  
)  
Number of learnable parameters' sets: 4  
torch.Size([30, 4])  
torch.Size([30])  
torch.Size([3, 30])  
torch.Size([3])
```


Epoch 1000/1000, Train Loss: 0.4989, Test Loss: 0.4967
Accuracy: 83.3%

Weights data have been saved in best_weights_SGD.pth



Epoch 1000/1000, Train Loss: 0.1113, Test Loss: 0.2206
Accuracy: 86.7%

Weights data have been saved in best_weights_Adam.pth




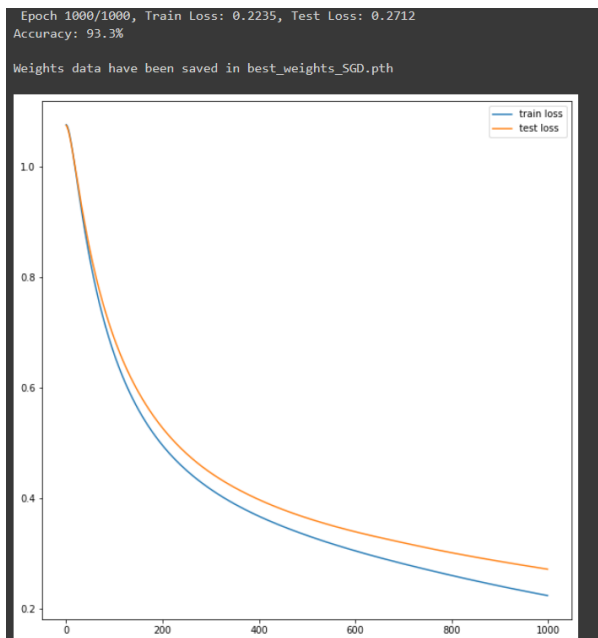
Precision, Recall, Confusion matrix, in testing

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 11 |
| 1 | 0.667 | 1.000 | 0.800 | 8 |
| 2 | 1.000 | 0.636 | 0.778 | 11 |
| accuracy | | | 0.867 | 30 |
| macro avg | 0.889 | 0.879 | 0.859 | 30 |
| weighted avg | 0.911 | 0.867 | 0.865 | 30 |

```
[[11 0 0]  
 [ 0 8 0]  
 [ 0 4 7]]
```

NN with 4:30:3 with Relu activation

```
NN_ReLU(  
  (flatten): Flatten(start_dim=1, end_dim=-1)  
  (linear_stack): Sequential(  
    (0): Linear(in_features=4, out_features=30, bias=True)  
    (1): ReLU()  
    (2): Linear(in_features=30, out_features=3, bias=True)  
  )  
)  
Number of learnable parameters' sets: 4  
torch.Size([30, 4])  
torch.Size([30])  
torch.Size([3, 30])  
torch.Size([3])
```



Confusion map and statistics for best outcome

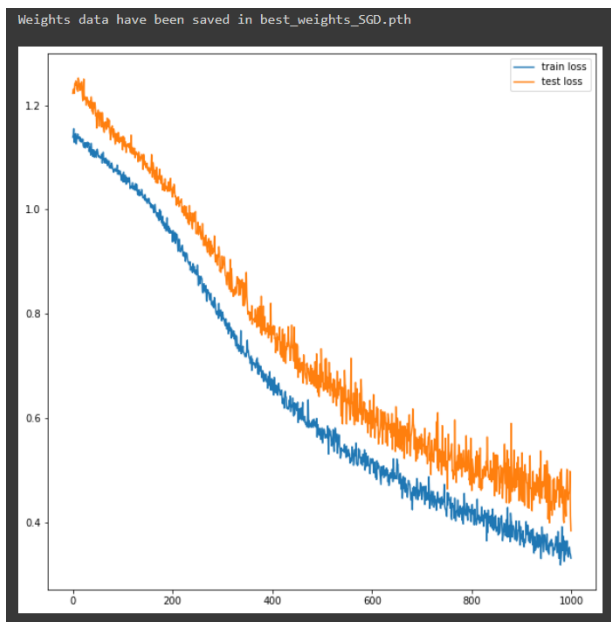
```
Precision, Recall, Confusion matrix, in testing
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 5 |
| 1 | 0.929 | 1.000 | 0.963 | 13 |
| 2 | 1.000 | 0.917 | 0.957 | 12 |
| accuracy | | | 0.967 | 30 |
| macro avg | 0.976 | 0.972 | 0.973 | 30 |
| weighted avg | 0.969 | 0.967 | 0.967 | 30 |

```
[[ 5  0  0]  
 [ 0 13  0]  
 [ 0  1 11]]
```

My NN 4:20:10:3 with Relu

```
My_NN_RelU(
  (linear_stack): Sequential(
    (0): Linear(in_features=4, out_features=20, bias=True)
    (1): ReLU()
    (2): Linear(in_features=20, out_features=10, bias=True)
    (3): ReLU()
    (4): Linear(in_features=10, out_features=3, bias=True)
  )
  (dropout): Dropout(p=0.1, inplace=False)
)
Number of learnable parameters' sets: 6
torch.Size([20, 4])
torch.Size([20])
torch.Size([10, 20])
torch.Size([10])
torch.Size([3, 10])
torch.Size([3])
```



Precision, Recall, Confusion matrix, in training

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 43 |
| 1 | 0.846 | 0.805 | 0.825 | 41 |
| 2 | 0.789 | 0.833 | 0.811 | 36 |
| accuracy | | | 0.883 | 120 |
| macro avg | 0.879 | 0.879 | 0.879 | 120 |
| weighted avg | 0.884 | 0.883 | 0.883 | 120 |


```
[[43 0 0]
 [ 0 33 8]
 [ 0 6 30]]
```

Precision, Recall, Confusion matrix, in testing

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 7 |
| 1 | 0.800 | 0.889 | 0.842 | 9 |
| 2 | 0.923 | 0.857 | 0.889 | 14 |
| accuracy | | | 0.900 | 30 |
| macro avg | 0.908 | 0.915 | 0.910 | 30 |
| weighted avg | 0.904 | 0.900 | 0.901 | 30 |

```
[[ 7 0 0]
 [ 0 8 1]
 [ 0 2 12]]
```

Add Learning rate Scheduler



Precision, Recall, Confusion matrix, in training

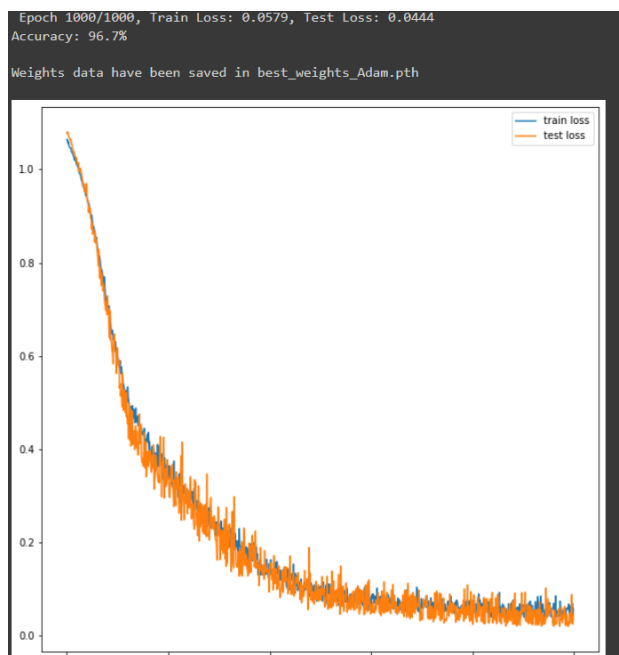
| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.933 | 1.000 | 0.966 | 42 |
| 1 | 1.000 | 0.949 | 0.974 | 39 |
| 2 | 1.000 | 0.974 | 0.987 | 39 |
| accuracy | | | 0.975 | 120 |
| macro avg | 0.978 | 0.974 | 0.975 | 120 |
| weighted avg | 0.977 | 0.975 | 0.975 | 120 |

```
[[42 0 0]
 [ 2 37 0]
 [ 1 0 38]]
```

Precision, Recall, Confusion matrix, in testing

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 8 |
| 1 | 0.846 | 1.000 | 0.917 | 11 |
| 2 | 1.000 | 0.818 | 0.900 | 11 |
| accuracy | | | 0.933 | 30 |
| macro avg | 0.949 | 0.939 | 0.939 | 30 |
| weighted avg | 0.944 | 0.933 | 0.933 | 30 |

```
[[ 8 0 0]
 [ 0 11 0]
 [ 0 2 9]]
```



Precision, Recall, Confusion matrix, in training

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.976 | 1.000 | 0.988 | 41 |
| 1 | 0.977 | 0.977 | 0.977 | 43 |
| 2 | 0.971 | 0.944 | 0.958 | 36 |
| accuracy | | | 0.975 | 120 |
| macro avg | 0.975 | 0.974 | 0.974 | 120 |
| weighted avg | 0.975 | 0.975 | 0.975 | 120 |


```
[[41 0 0]
 [ 0 42 1]
 [ 1 1 34]]
```

Precision, Recall, Confusion matrix, in testing

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 9 |
| 1 | 1.000 | 1.000 | 1.000 | 7 |
| 2 | 1.000 | 1.000 | 1.000 | 14 |
| accuracy | | | 1.000 | 30 |
| macro avg | 1.000 | 1.000 | 1.000 | 30 |
| weighted avg | 1.000 | 1.000 | 1.000 | 30 |

```
[[ 9 0 0]
 [ 0 7 0]
 [ 0 0 14]]
```

Adam with wreigth decay



Precision, Recall, Confusion matrix, in training

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.956 | 1.000 | 0.977 | 43 |
| 1 | 0.909 | 0.750 | 0.822 | 40 |
| 2 | 0.786 | 0.892 | 0.835 | 37 |
| accuracy | | | 0.883 | 120 |
| macro avg | 0.883 | 0.881 | 0.878 | 120 |
| weighted avg | 0.888 | 0.883 | 0.882 | 120 |

```
[[43 0 0]
 [ 1 30 9]
 [ 1 3 33]]
```

Precision, Recall, Confusion matrix, in testing

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 7 |
| 1 | 1.000 | 0.800 | 0.889 | 10 |
| 2 | 0.867 | 1.000 | 0.929 | 13 |
| accuracy | | | 0.933 | 30 |
| macro avg | 0.956 | 0.933 | 0.939 | 30 |
| weighted avg | 0.942 | 0.933 | 0.932 | 30 |

```
[[ 7 0 0]
 [ 0 8 2]
 [ 0 0 13]]
```

Scaler :

We introduced a standard deviation scaler to standardize features of the data set by scaling to unit variance and removing the mean using column summary statistics on the samples in the training set. This method helps with data far away from center that otherwise would contribute very small amount to the change.

Sigmoid vs Relu :

Because sigmoid compresses back propagations adjustments needs more repetitions-epoch to reach a optimum solution. If we let epochs be bigger the network would give a good solution because the data are very small and the network has one hidden layer. Although the data and the networks layout there is a difference with Relu activation function producing better results in average.

Std "with momentum" :

By introducing a momentum in optimizer we overcome the small "hills" by not following the gradient all the time like a heavy ball.

Scheduler :

Scheduler slowly decreases learning rate when the system does not move a lot. With that addition we can "reach deeper" in lower minimums sometimes. If the learning rate is too high the step accordingly would be too long and we would not reach the true minimum.


Std vs Adam :

Adam method introduces two optimisations :

- Adaptive Gradient Algorithm (AdaGrad) that maintains a per-parameter learning rate that improves performance on problems with sparse gradients (e.g. natural language and computer vision problems).
- Root Mean Square Propagation (RMSProp) that also maintains per-parameter learning rates that are adapted based on the average of recent magnitudes of the gradients for the weight (e.g. how quickly it is changing). This means the algorithm does well on online and non-stationary problems (e.g. noisy).

Source: [Gentle Introduction to the Adam Optimization Algorithm for Deep Learning - MachineLearningMastery.com](https://machinelearningmastery.com/gentle-introduction-to-the-adam-optimization-algorithm-for-deep-learning/)

Without all the optimizations on NN Sigmoid 30 :




Precision, Recall, Confusion matrix, in training

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.000 | 0.000 | 0.000 | 35 |
| 1 | 0.000 | 0.000 | 0.000 | 42 |
| 2 | 0.361 | 1.000 | 0.531 | 43 |
| accuracy | | | 0.358 | 120 |
| macro avg | 0.120 | 0.333 | 0.177 | 120 |
| weighted avg | 0.129 | 0.358 | 0.190 | 120 |

```
[[ 0 1 34]
 [ 0 0 42]
 [ 0 0 43]]
```

After epoch 10000



Precision, Recall, Confusion matrix, in training

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 38 |
| 1 | 1.000 | 0.475 | 0.644 | 40 |
| 2 | 0.667 | 1.000 | 0.800 | 42 |
| accuracy | | | 0.825 | 120 |
| macro avg | 0.889 | 0.825 | 0.815 | 120 |
| weighted avg | 0.883 | 0.825 | 0.811 | 120 |

```
[[38 0 0]
 [ 0 19 21]
 [ 0 0 42]]
```

Precision, Recall, Confusion matrix, in testing

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.000 | 1.000 | 1.000 | 12 |
| 1 | 1.000 | 0.600 | 0.750 | 10 |
| 2 | 0.667 | 1.000 | 0.800 | 8 |
| accuracy | | | 0.867 | 30 |
| macro avg | 0.889 | 0.867 | 0.850 | 30 |
| weighted avg | 0.911 | 0.867 | 0.863 | 30 |

```
[[12 0 0]
 [ 0 6 4]
 [ 0 0 8]]
```

With 10x times the epoch the network cant even reach the worst case of the nectorks above